

Hellomatik

Checking funding proposals against the call rules, with every finding traced to its source

Catalyst Grant 2026 application · HELLOMATIK, S.L. (Madrid, Spain) · Contact: Isaac Correa Villa, isaac.correa@hellomatik.com · hellomatik.com

1. The problem

Every competitive funding call publishes rules that a proposal must meet: eligibility conditions, page and format limits, mandatory sections, budget categories and rates, and declarations. Applicants, research offices and funder staff check these rules by hand, usually under deadline pressure. Proposals are declared inadmissible or lose points for errors that a careful reading would have caught: a budget table that does not match the narrative, a page limit exceeded, a person-month total that differs between two parts, a claim in the CV that no source supports.

It affects research-office and pre-award staff, who check many proposals per call cycle; applicants without a grants office; and funder staff who screen admissibility. The cost is staff hours at the busiest point of the year and, when an error slips through, weeks of work losing points or being rejected on form.

We met this problem ourselves. In September 2026 we prepared several European proposals: an EIC Pathfinder Challenge proposal of 17 pages with an official budget workbook, an EIC Accelerator short application, and EuroHPC access applications. Checking them against the call documents took more time than writing some sections.

2. Your workflow

Trigger: an applicant or grants officer uploads a draft proposal (Part A form export, Part B, budget workbook, annexes) and selects the call.

1. **Read the call.** The agent extracts the rules from the call documents (work programme topic, template, admissibility and eligibility conditions, budget rules) into a written list, each rule linked to the clause it came from. A grants officer confirms the list once per call; later proposals reuse it.
2. **Map the proposal.** It extracts the checkable facts from the proposal: page counts, section headings, budget figures per category and work package, person-months, deliverable dates, named partners, and factual claims about the team and prior results.
3. **Check deterministically where possible.** Arithmetic (totals, rates, indirect costs), limits and cross-part consistency are computed by code, not by the language model.
4. **Check claims against sources.** For factual claims the agent looks for supporting evidence in the documents the applicant provides (CV, publication records with DOIs, registry extracts) and marks each claim as supported, contradicted or unsupported.
5. **Report findings.** Each finding states the rule, the proposal passage, the evidence and a suggested correction.
6. **Human decision.** The applicant or grants officer accepts, edits or dismisses each finding; nothing is changed in the proposal automatically. Decisions are recorded.

Where it runs: first as a standalone review step used by applicants and research offices before submission; the target is to sit inside the pre-award process that research offices already run, next to the research-information and grants systems they use every day.

3. Trust, audit and governance

This is the part Hellomatik already runs in production for businesses. Our platform writes down how an organisation decides (its data, its rules, the actions allowed and who approves each one) and runs AI agents only through those named rules. Every figure an agent reports is computed by the same function that feeds the user's screen, never by the language model, and agents have read-only access; named actions that wait for a person's approval are the next layer of the platform.

- **What the agent did and why:** each finding shows the rule and its clause in the call document, the exact proposal passage, the evidence used and whether the check was computed or model-judged.
- **Provenance:** rules cite call clauses; facts cite page and paragraph of the proposal; evidence cites the source document or DOI. A finding without a citation is not shown.
- **Uncertainty:** model-judged findings carry a confidence label and are never presented as computed results. When the call text is ambiguous, the agent escalates the question to the grants officer instead of deciding.
- **Accountability:** the person who accepts or dismisses a finding is recorded; the applicant remains responsible for the proposal, as funders such as the European Commission require.

4. Team

Isaac Correa Villa, co-founder and joint administrator, leads product, engineering and research. He is the author of the technical report *How we build a company's ontology* (2026, DOI 10.5281/zenodo.22647735) and of *The machine of the future* (DOI 10.5281/zenodo.22671771); ORCID 0009-0003-0657-4419. Before Hellomatik he co-founded Kodalogic and delivered about 300 business dashboards to clients in the United States, Europe, South America and South Korea, with 103 verified client reviews.

Daniel Correa Villa, co-founder and joint administrator.

We picked this problem because our core technology is rule extraction, deterministic checking and human approval, and because we have just lived it as applicants to European calls.

5. Where you are today

Stage: the underlying platform is a working product in use with business customers in four industries, on our own servers in the EU. The proposal-checking workflow is a **concept** built on it; the rule extraction, deterministic checks, citations and human approval it needs already run in production for business decisions. What follows is what broke in that production work and what we changed, because the same failure modes apply to funding proposals.

A language model invented the rules. For a multinational B2B manufacturer, a first model of the business (objects, keys, statuses, actions) was drafted with language-model assistance from documentation and screens. It named the core objects correctly, but it invented properties, primary keys and a status list that do not exist (the real field is a one-character code), invented an after-sales object that the real system models differently, left out the whole financial, logistics and after-sales domain, and mixed planned features with existing ones. We discarded it and rebuilt the model by reading 84 database migrations and 36 service endpoints, citing every element to its source.

Consequence for this project: rules extracted from a call document are never used until each one cites a clause and a person confirms the list.

Totals matched to the cent, and still there was a gap. Integrating an advertising account, spend reconciled exactly across daily, campaign and advert level, but differed by 7.88 EUR from the account's lifetime total. The cause was the source's 37-month data retention window: older data could no longer be retrieved. The agent now states such gaps and their cause instead of hiding or "fixing"

them. *Consequence*: a budget check that cannot reconcile reports the difference and the likely cause, and never adjusts a figure.

Tested questions it must refuse or admit. The agents run against a written list of trap questions with expected answers: a return-on-ad-spend question when no purchases were recorded ("no purchase conversions are recorded"), a year with no data, a forecast request, a non-existent identifier, and write requests (pause a campaign, raise a budget) that must be refused because the access is read-only. *Consequence*: the proposal checker gets the same kind of trap suite (a rule the call does not contain, a claim with no source, a figure it cannot compute) before any pilot.

Routing to the wrong source. In workspaces with several connected sources, a question could be routed to the wrong one unless the user named the source. *Consequence*: every finding names the document it read, and an ambiguous question is escalated instead of answered.

Not built yet: approval-gated business actions are the next layer; this checker would be one of its first uses.

6. Alternatives and competitors

Today it is solved by hand, with checklists from national contact points and consultancies, or not at all. AI grant-writing tools focus on drafting text and offer compliance checks as a feature, for example EUACC (euacc.ai), GrantForge (grantforge.eu) and Granted (grantedai.com). We differ in three ways: rules are extracted once, confirmed by a person and cited to the call clause; arithmetic and consistency are computed deterministically rather than judged by a language model; and every finding requires a human decision that is recorded, which matters now that funders make applicants fully responsible for AI-assisted content.

7. Where this goes

A review layer for the whole pre-award cycle: applicants check drafts, research offices check many proposals against one confirmed rule set, and funders screen admissibility with the same cited findings. Research offices would pay per seat or call cycle; funders would license it. Each call is modelled once and reused.

8. Fit with Digital Science

We sit in **funding decisions**, for institutions and funders. Digital Science works with research offices and funders every day, and its research information and grants data would let the agent check claims about prior funding and outputs against trusted records. We would like to build this with the people who know how research offices work.

9. Budget

£25,000 would fund three months of engineering to turn the concept into an open prototype on real Horizon Europe templates (£17,000), a measured pilot with one research office on a live call, recording findings accepted, false alarms and staff minutes saved (£6,000), and model and hosting costs (£2,000). Without it, this work competes with customer projects and waits.