

Hellomatik

Checking funding proposals against the call rules, with every finding traced to its source

Catalyst Grant 2026 application · HELLOMATIK, S.L. (Madrid, Spain) · Contact: Isaac Correa Villa, isaac.correa@hellomatik.com · hellomatik.com

1. The problem

Every competitive funding call publishes rules that a proposal must meet: eligibility conditions, page and format limits, mandatory sections, budget categories and rates, and declarations. Applicants, research offices and funder staff check these rules by hand, usually under deadline pressure. Proposals are declared inadmissible or lose points for errors that a careful reading would have caught: a budget table that does not match the narrative, a page limit exceeded, a person-month total that differs between two parts, a claim in the CV that no source supports.

The people affected are research-office and pre-award staff, who check many proposals per call cycle; applicants, especially small teams and SMEs without a grants office; and funder programme staff, who screen admissibility before evaluation. The cost is staff hours at the busiest point of the year and, when an error slips through, a proposal that took weeks of work losing points or being rejected on form.

We met this problem ourselves. In September 2026 we prepared several European proposals: an EIC Pathfinder Challenge proposal of 17 pages with an official budget workbook, an EIC Accelerator short application, and EuroHPC access applications. Checking them against the call documents took more time than writing some sections, and it still found real errors late (examples in section 5).

2. Your workflow

Trigger: an applicant or grants officer uploads a draft proposal (Part A form export, Part B, budget workbook, annexes) and selects the call.

1. **Read the call.** The agent extracts the rules from the call documents (work programme topic, template, admissibility and eligibility conditions, budget rules) into a written list, each rule linked to the clause it came from. A grants officer confirms the list once per call; later proposals reuse it.
2. **Map the proposal.** It extracts the checkable facts from the proposal: page counts, section headings, budget figures per category and work package, person-months, deliverable dates, named partners, and factual claims about the team and prior results.
3. **Check deterministically where possible.** Arithmetic (totals, rates, indirect costs), limits and cross-part consistency are computed by code, not by the language model.
4. **Check claims against sources.** For factual claims the agent looks for supporting evidence in the documents the applicant provides (CV, publication records with DOIs, registry extracts) and marks each claim as supported, contradicted or unsupported.
5. **Report findings.** Each finding states the rule, the proposal passage, the evidence and a suggested correction.

6. **Human decision.** The applicant or grants officer accepts, edits or dismisses each finding; nothing is changed in the proposal automatically. Decisions are recorded.

Where it runs: first as a standalone review step used by applicants and research offices before submission; the target is to sit inside the pre-award process that research offices already run, next to the research-information and grants systems they use every day.

3. Trust, audit and governance

This is the part Hellomatik already runs in production for businesses. Our platform writes down how an organisation decides (its data, its rules, the actions allowed and who approves each one) and runs AI agents only through those named rules. Every figure an agent reports is computed by a versioned formula over source data, and every action with an external effect waits for a person's approval, with the decision recorded.

- **What the agent did and why:** each finding shows the rule and its clause in the call document, the exact proposal passage, the evidence used and whether the check was computed or model-judged.
- **Provenance:** rules cite call clauses; facts cite page and paragraph of the proposal; evidence cites the source document or DOI. A finding without a citation is not shown.
- **Uncertainty:** model-judged findings carry a confidence label and are never presented as computed results. When the call text is ambiguous, the agent escalates the question to the grants officer instead of deciding.
- **Accountability:** the person who accepts or dismisses a finding is recorded; the applicant remains responsible for the proposal, as funders such as the European Commission require.

4. Team

Isaac Correa Villa, co-founder and joint administrator, leads product, engineering and research. He is the author of the technical report *How we build a company's ontology* (2026, DOI 10.5281/zenodo.22647735) and of *The machine of the future* (DOI 10.5281/zenodo.22671771); ORCID 0009-0003-0657-4419. Before Hellomatik he co-founded Kodalogic and delivered about 300 business dashboards to clients in the United States, Europe, South America and South Korea, with 103 verified client reviews.

Daniel Correa Villa, co-founder and joint administrator.

We picked this problem because our core technology is rule extraction, deterministic checking and human approval with a record, and because we have just lived the problem as applicants to European calls.

5. Where you are today

Stage: the underlying platform is a working product in use with business customers in four industries, running on our own servers in the EU. The proposal-checking workflow is at **concept stage with a manual trial:** we applied the steps above, with AI assistance and human review, to our own September 2026 proposals.

What the trial found (real findings, corrected in every document that was still editable):

- A team claim ("clients in 27 countries") in a Part B and a pitch deck that did not match the approved, verifiable biography; flagged as unsupported and replaced in the draft still open.
- An inconsistent count between documents ("six" EuroHPC applications in one, five submitted in reality).
- A third-party funding guide stating that the EIC Pathfinder Challenges require three partners; checked against the official Work Programme, where single legal entities are eligible. The agent must escalate conflicting sources rather than trust the first one.
- Budget workbook totals per work package reconciled against the Part B tables (1,835,625 EUR in both).

Where it broke: checks on the funder's online forms could not be automated when the form itself failed (the portal's cost calculator returned server errors), which shows the agent must report "could not check" explicitly instead of passing silently.

6. Alternatives and competitors

Today the problem is solved by hand, with checklists from national contact points and consultancies, or not at all. AI grant-writing tools focus on drafting text and offer compliance checks as a feature, for example EUACC (euacc.ai), GrantForge (grantforge.eu) and Granted (grantedai.com). We differ in three ways: rules are extracted once, confirmed by a person and cited to the call clause; arithmetic and consistency are computed deterministically rather than judged by a language model; and every finding requires a human decision that is recorded, which matters now that funders make applicants fully responsible for AI-assisted content.

7. Where this goes

A review layer for the whole pre-award cycle: applicants check drafts, research offices check many proposals against the same confirmed rule set, and funders screen admissibility with the same cited findings. Research offices and consultancies would pay per call cycle or per seat; funders would license it for admissibility screening. The rule sets per call can be shared, so each call is modelled once.

8. Fit with Digital Science

We sit in **funding decisions**, for institutions and funders. Digital Science works with research offices and funders every day, and its research information and grants data would let the agent check claims about prior funding and outputs against trusted records. We would like to build this with the people who know how research offices work.

9. Budget

£25,000 would fund three months of engineering to turn the manual trial into an open prototype on real Horizon Europe templates (£17,000), a measured pilot with one research office on a live call, recording findings accepted, false alarms and staff minutes saved (£6,000), and model and hosting costs (£2,000). Without it, this work competes with customer projects and waits.